


Modeling Interactions with Deep Learning

Jack Lanchantin | PhD Thesis Defense | July 20th 2021


University of Virginia Department of Computer Science

Committee: Vicente Ordoñez (chair), Yanjun Qi (advisor), Yangfeng Ji, Clint Miller, Casey Greene


tree bird bench dog




the cat sat on the small mat.




ATCGATCGATGCCGTTACGATGATCGTACGAAT



Machine learning models




Machine learning models with relational biases




Thesis hypothesis

Deep learning with relational biases can be used to understand and improve tasks that involve complex interactions.


Using Deep Learning to Model Interactions




interactions between
input features



interactions between
input samples



interactions between
output labels



Genomic Sequences

ChromeGCN
Lanchantin et al.
ECCB 2020

Deep Motif Dashboard
Lanchantin et al.
PSB 2017

Memory Matching Networks
Lanchantin et al.
ICLR Workshops 2017

DeepChrome
Singh et al.
ECCB 2016

AttentiveChrome
Singh et al.
NeurIPS 2017

Proteins

DeepVHPII
Lanchantin et al.
ACM-BCB 2021

MUST-CNN
Lin et al.
AAAI 2016

Visual Objects & Natural Language

C-Tran
Lanchantin et al.
CVPR 2021

LaMP
Lanchantin et al.
ECML 2019

Re-evaluating Adversarial
Morris et al.
EMNLP Findings 2020

Deep WordBug
Gao et al.
DLS 2018

Genomic Sequences

ChromeGCN
Lanchantin et al.
ECCB 2020

Deep Motif Dashboard
Lanchantin et al.
PSB 2017

Memory Matching Networks
Lanchantin et al.
ICLR Workshops 2017

DeepChrome
Singh et al.
ECCB 2016

AttentiveChrome
Singh et al.
NeurIPS 2017

Proteins

DeepVHPI
Lanchantin et al.
ACM-BCB 2021

MUST-CNN
Lin et al.
AAAI 2016

Visual Objects & Natural Language

C-Tran
Lanchantin et al.
CVPR 2021


LaMP
Lanchantin et al.
ECML 2019

Re-evaluating Adversarial
Morris et al.
EMNLP Findings 2020


Deep WordBug
Gao et al.
DLS 2018

Modeling Genomic Sequence Interactions with ChromeGCN


interactions between
input features




interactions between
input samples



interactions between
output labels



What is DNA?



What is DNA?

ATGCTCGATGCTAATA
CGACTGAGATTACTGAGACTGAGACTCTAGAT

What is DNA?

ACTGCTACCTATGACGTGATGCATCGTAGCTA

← 250M chars →


chromosome 1

Chromatin Profile

Chromatin profile signals: elements that control gene expression
(transcription factors, histone modifications, DNA accessibility)



Local Sequence Chromatin Profile Prediction



Local Sequence Chromatin Profile Prediction

ACTGCTAC

x^1

← 1k chars →

CTATGACG

x^2


TGATGCAT

x^3

CGTAGCTA


x^4

Local Sequence Chromatin Profile Prediction




Zhou and Troyanskaya 2015, Alipanahi et al. 2015, Kelley et al. 2016


Influence of Long-Range Interactions on Chromatin Profile




Influence of Long-Range Interactions on Chromatin Profile



Genome: Locally a Sequence, Globally a Graph




Genome: Locally a Sequence, Globally a Graph



"Hi-C" interaction data




Hyperlinks



Research Question 1

Can we model both local sequence features and long range interactions for chromatin profile prediction?

Graph Neural Networks (GNNs)




$$G = (V, E)$$

$$\mathbf{v}_i^t \in \mathbb{R}^d$$

Scarselli et al. 2008, Li et al. 2015, Kipf et al. 2016, Gilmer et al. 2017, Bronstein et al. 2017, Battaglia et al. 2018

Graph Neural Networks (GNNs)




$$G = (V, E)$$
$$\mathbf{v}_i^t \in \mathbb{R}^d$$

message function F_m $\mathbf{m}_i^t = \sum_{j \in \mathcal{N}(i)} F_m(\mathbf{v}_i^t, \mathbf{v}_j^t),$

node update function F_u $\mathbf{v}_i^{t+1} = F_u(\mathbf{m}_i^t)$

Scarselli et al. 2008, Li et al. 2015, Kipf et al. 2016, Gilmer et al. 2017, Bronstein et al. 2017, Battaglia et al. 2018

Graph Neural Networks (GNNs)




$$G = (V, E)$$
$$\mathbf{v}_i^t \in \mathbb{R}^d$$

message function F_m $\mathbf{m}_i^t = \sum_{j \in \mathcal{N}(i)} F_m(\mathbf{v}_i^t, \mathbf{v}_j^t),$

node update function F_u $\mathbf{v}_i^{t+1} = F_u(\mathbf{m}_i^t)$

Scarselli et al. 2008, Li et al. 2015, Kipf et al. 2016, Gilmer et al. 2017, Bronstein et al. 2017, Battaglia et al. 2018


Graph Neural Networks (GNNs)




Graph neural networks are good at modeling *interactions*.

ChromeGCN: Combining Local Sequence and Long Range Interactions for Chromatin Profile Prediction

ChromeGCN




ChromeGCN




ChromeGCN

$$\mathcal{L} = \frac{1}{L} \sum_{i=1}^L -(y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$$




ChromeGCN




Chromatin Profile Prediction: Experimental Setup

Chromatin profile labels and interaction graph data: transcription factor, histone modification, DNA accessibility, Hi-C, for two cell lines (GM12878, K562)


Data splits: 16 training chromosomes, 3 validation, 3 testing



Chromatin Profile Prediction: Results



Chromatin Profile Prediction: Results




modeling long range interactions is important for signals with many interactions (high degree)

↑ ChromeGCN is better
↓ CNN is better

Implications of Predictive Genomic Sequence Models

1. Mutation Analysis




GATGCTGCTCTACGA
GATGCTG_{AT}CTACGA

2. Feature Attribution



Implications of Predictive Genomic Sequence Models

1. Mutation Analysis




GATGCTGCTCTACGA
GATGCTGATCTACGA


2. Feature Attribution



Deep Motif Dashboard: Visualizing Genomic Sequence Syntax




Sequence saliency maps: finding important sequence features




$$\hat{y}_+(X) \approx w^T X + b = \sum_{i=1}^{|X|} w_i x_i$$
$$w = \frac{\partial \hat{y}_+}{\partial X} \Big|_{X_0} = \text{"sequence saliency map"}$$

Sequence saliency maps: finding important sequence features




Sequence saliency maps: finding important sequence features




Local motif interactions are important!

Hi-C saliency maps: identifying important long range interactions




Contributions


1. Cohesive framework: we fuse local sequence features and long range interactions for chromatin profile prediction



2. Accurate: incorporating long range interactions outperforms the baselines



3. Interpretable: we introduce deep motif dashboard to interpret local features and Hi-C saliency maps to find important interactions



Graph neural network models can effectively exploit and discover important interactions in DNA for functional genomics

Genomic Sequences

ChromeGCN
Lanchantin et al.
ECCB 2020

Deep Motif Dashboard
Lanchantin et al.
PSB 2017

Memory Matching Networks
Lanchantin et al.
ICLR Workshops 2017

DeepChrome
Singh et al.
ECCB 2016

AttentiveChrome
Singh et al.
NeurIPS 2017

Proteins

DeepVHPPPI
Lanchantin et al.
ACM-BCB 2021

MUST-CNN
Lin et al.
AAAI 2016

Visual Objects & Natural Language

C-Tran
Lanchantin et al.
CVPR 2021


LaMP
Lanchantin et al.
ECML 2019

Re-evaluating Adversarial
Morris et al.
EMNLP Findings 2020


Deep WordBug
Gao et al.
DLS 2018

Modeling Host-Virus Protein-Protein Interactions with DeepVHPI


interactions between
input features




interactions between
input samples




interactions between
output labels





What is a protein?




sequences with alphabet of 20 characters

What is a protein?

MQGHFTETKHE

sequences with alphabet of 20 characters

What is a protein?




3D Structure



MHFTEDKATILWGKVNVGETLGRVYPWQ

Primary Sequence

What is a protein?



3D Structure




MHFTEDKATILWGKVNVGETLGRVYPWQ

Primary Sequence


Structure Determines Function




Protein-Protein Interactions (PPIs)




Virus-Host Protein-Protein Interactions



Novel Virus-Human Protein Interaction Prediction from Sequence




Novel Virus-Human Protein Interaction Prediction from Sequence



Gomez et al. 2003, Ben-Hur and Noble 2005, Qi et al. 2010, Eid et al. 2015, Yang et al. 2020

Novel Virus-Human Protein Interaction Prediction from Sequence




structure (3D shape) is important, but hard to obtain

Research Question 2


Can we incorporate individual protein structure into a sequence based protein-protein interaction prediction model?

Structure Prediction from Sequence




Lin, Lanchantin, Qi 2016, Rives et al. 2019, Rao et al. 2020, Jumper et al. 2020

Transfer Learning




Glorot *et al.* 2011, Devlin *et al.* 2018


Transfer Learning for Sequence-Based Interaction Prediction




Transfer Learning for Sequence-Based Interaction Prediction




Masked Language Model (MLM) Pretraining




Transfer Learning for Sequence-Based Interaction Prediction




Structure Prediction (SP) Pretraining




Transfer Learning for Sequence-Based Interaction Prediction




Interaction Prediction




Interaction Prediction



Transfer Learning for Sequence-Based Interaction Prediction




Transfer Learning for Sequence-Based Interaction Prediction




Transformers

Word relationships are learned via neural attention



Vaswani et al. 2017, Devlin et al. 2018

Neural Attention



Attention coefficient e_{ij}^t for pair of nodes $(\mathbf{v}_i^t, \mathbf{v}_j^t)$


$$e_{ij}^t = (\mathbf{W}^q \mathbf{v}_i^t)^\top (\mathbf{W}^u \mathbf{v}_j^t)$$

Attention weight α_{ij}^t produced by normalizing over neighboring nodes

$$\alpha_{ij}^t = \text{softmax}_j(e_{ij}^t) = \frac{\exp(e_{ij}^t)}{\sum_{k \in \mathcal{N}(i)} \exp(e_{ik}^t)}.$$


Bahdanau et al. 2013, Vaswani et al. 2017

Neural Attention




Attention allows us to learn interaction
importance when it is not known or predefined

Transformers for Protein Sequences




Transformers for Protein Sequences




Rives et al. 2019, Rao et al. 2020

Transformers for Protein Sequences




protein motifs (i.e. “words”) influence protein-protein interactions


DeepVHPI Transformer




DeepVHPI Transformer



DeepVHPII Transformer




DeepVHPII Transformer



Use Cases of Sequence Based Interaction Predictors

1. predict novel virus interactions




2. analyze how mutations affect interactions




Use Cases of Sequence Based Interaction Predictors

1. predict novel virus interactions



2. analyze how mutations affect interactions




Protein-Protein Interaction Prediction: Experimental Setup

Training Data: HPIDB 3.0 Dataset

- 22,000 positive interactions, 226,000 negative interactions
- 1,100 virus proteins, 20,000 host (human) proteins

Testing Data:

- H1N1, Ebola, SARS-CoV-2 proteins



Protein-Protein Interaction Prediction: Results


Method	H1N1 ¹		Ebola ¹		SARS-CoV-2 ²	
	AUROC	F1	AUROC	F1	AUROC	F1
SVM (Zhou et al.)	0.886	0.762	0.867	0.760	-	-
Embedding + RF (Yang et al)	-	-	-	-	0.748	0.115
DeepVHPPI	0.906	0.821	0.927	0.836	0.726	0.089
DeepVHPPI + MLM	0.920	0.804	0.943	0.867	0.735	0.095
DeepVHPPI + MLM + SP	0.957	0.867	0.976	0.895	0.767	0.105

¹Even positive/negative testing split ²Imbalanced positive/negative testing split

Transformers and transfer learning helps in all three testing datasets

Use Cases of Sequence Based Interaction Predictors

1. predict novel virus interactions



2. analyze how mutations affect interactions



Perturbation Analysis: Investigating Mutations

Short Article

D614G Spike Mutation Increases SARS CoV-2 Susceptibility to Neutralization

Drew Weiss¹,
Hornsby²,
⁷, Katayoun³,
Lin⁹, Ying¹

The NEW ENGLAND JOURNAL of MEDICINE

CLINICAL IMPLICATIONS OF BASIC RESEARCH

Elizabeth G. Phimister, Ph.D., Editor


Emergence of a Highly Fit SARS-CoV-2 Variant

nature


<https://doi.org/10.1038/s41586-021-03777-9>

Accelerated Article Preview


Reduced sensitivity of SARS-CoV-2 variant Delta to antibody neutralization




Perturbation Analysis: Investigating Mutations




Perturbation Analysis: Investigating Mutations




Perturbation Analysis: Investigating Mutations




Perturbation Analysis



Perturbation Analysis




Perturbation Analysis




Experimental Setup

105,528 mutated Spike sequences and corresponding ACE2 binding affinities from Starr et al.



Perturbation Analysis: Mutated Spike and ACE2 Interactions


Correlation between DeepVHPPI binding prediction and dissociation constant




DeepVHPPI can be used to help understand the **effects of virus variants**

Contributions


1. Flexible transfer learning framework for protein-protein interaction prediction



2. Accurate novel virus interaction predictions



3. Interpretable and interactive mutation perturbation analysis



Transfer learning and Transformers can be used to accurately predict protein-protein interactions

Genomic Sequences

ChromeGCN
Lanchantin et al.
ECCB 2020

Deep Motif Dashboard
Lanchantin et al.
PSB 2017

Memory Matching Networks
Lanchantin et al.
ICLR Workshops 2017

DeepChrome
Singh et al.
ECCB 2016

AttentiveChrome
Singh et al.
NeurIPS 2017

Proteins

DeepVHPII
Lanchantin et al.
ACM-BCB 2021

MUST-CNN
Lin et al.
AAAI 2016

Visual Objects & Natural Language

C-Tran
Lanchantin et al.
CVPR 2021


LaMP
Lanchantin et al.
ECML 2019

Re-evaluating Adversarial
Morris et al.
EMNLP Findings 2020


Deep WordBug
Gao et al.
DLS 2018

Modeling Image Label Interactions with C-Tran

interactions between
input features



interactions between
input samples



interactions between
output labels





Image Scene Understanding




Multi-label Classification

Given input x , predict the **set** of target labels $\{y_1, y_2, \dots, y_L\}$, $y_i \in \{0, 1\}$




Multi-label Classification

Given input x , predict the **set** of target labels $\{y_1, y_2, \dots, y_L\}$, $y_i \in \{0, 1\}$



Regular inference



Inference with partial knowledge

Information (e.g. certain labels) can be known **prior to** performing any **visual recognition**


x



raincoat = 1

Inference with partial knowledge

Information (e.g. certain labels) can be known **prior to** performing any **visual recognition**



Inference with context-specific information: a realistic setting

Geo-Location Tags



📍 Lake George, NY

lake

Social Media Tags



#avocado #bowl

avocado

bowl

News Captions




Mt. Sinabung sent smoke over western Indonesia Wednesday.

mountain

smoke

Inference with context-specific information: a realistic setting



problem: for each test image, known labels can be different

Inference with context-specific information: a realistic setting



#cup #table

cup

table


problem: for each test image, known labels can be different

Inference with context-specific information: a realistic setting



problem: for each test image, known labels can be different

Previous Work: Feedbackprop




Wang et al. 2018


Research Question 3

Can we model interactions between object labels, and also exploit the interactions from arbitrary amounts of known labels?


Classification Transformer (C-Tran)




Classification Transformer (C-Tran)




Classification Transformer (C-Tran)




Classification Transformer (C-Tran)




Classification Transformer (C-Tran)




Classification Transformer (C-Tran)




Classification Transformer (C-Tran)




Classification Transformer (C-Tran)




Classification Transformer (C-Tran)




Classification Transformer (C-Tran)




Classification Transformer (C-Tran)




Classification Transformer (C-Tran)



Classification Transformer (C-Tran)




Classification Transformer (C-Tran)



Three different inference settings

1. Regular Inference

predict masked labels



C-Tran




mask all labels

2. Partial Label Inference

predict masked labels



C-Tran




mask unknown labels

3. Extra Label Inference

predict masked labels




C-Tran




mask unknown labels

Three different inference settings


1. Regular Inference



2. Partial Label Inference



3. Extra Label Inference



Regular inference setting



0.8 improvement in mAP and F1 score




1.0 improvement in mAP and F1 score


Attention Visualization




Attention Visualization




Dependencies between labels and image patches



frisbee-to-image attention




Dependencies between labels and other labels




Three different inference settings


1. Regular Inference



2. Partial Label Inference



3. Extra Label Inference



Partial label inference setting

4 different percentages of known labels

Partial Labels Known	COCO-80	VG-500	NEWS-500	COCO-1000
	0% 25% 50% 75%	0% 25% 50% 75%	0% 25% 50% 75%	0% 25% 50% 75%



Partial label inference setting

4 different percentages of known labels



Partial Labels Known	COCO-80				VG-500				NEWS-500				COCO-1000			
	0%	25%	50%	75%	0%	25%	50%	75%	0%	25%	50%	75%	0%	25%	50%	75%
Feedbackprop [49]	80.1	80.6	80.8	80.9	29.6	30.1	30.8	31.6	14.7	21.1	23.7	25.9	29.2	30.1	31.5	33.0
C-Tran	85.1	85.2	85.6	86.0	38.4	39.3	40.4	41.5	18.1	29.7	35.5	39.4	34.3	35.9	37.4	39.1

C-Tran: improved accuracy with **any amount** of known labels


Partial label inference setting

Images	True Labels	ResNet-101	C-Tran	C-Tran + partial labels
 ID:000000362831	fork knife, spoon, bowl, chair, diningtable	fork, sandwich, diningtable, spoon, cup	fork, knife, diningtable, person, cake	spoon=1, trafficlight=0, bench=0, dog=0, ... fork, knife, diningtable, person, bowl


Three different inference settings

1. Regular Inference

predict masked labels



C-Tran



mask all labels

2. Partial Label Inference

predict masked labels



C-Tran




mask unknown labels

3. Extra Label Inference

predict masked labels




C-Tran



mask unknown labels

Extra label inference setting



Extra label inference setting on CUB bird classification dataset

4 different percentages of known extra labels



Extra Label Groups Known (ϵ)	0%	36%	54%	71%
Standard [23]	82.7	82.7	82.7	82.7
Multi-task [23]	83.8	83.8	83.8	83.8
ConceptBottleneck [23]	80.1	87.0	93.0	97.5
C-Tran	83.8	90.0	97.0	98.0

C-Tran: improved accuracy with **any amount** of extra attribute labels

Extra label inference setting



Image

True
Label

Anna
Hummingbird

C-Tran

Rufous
Hummingbird (96%)

C-Tran + Extra Labels


has_bill_shape_needle = 1,
has_wing_color_green=1,
has_upperparts_color=green=1,
has_back_color_blue=0,
has_back_color_brown=0

...

Anna
Hummingbird (99%)


Model intervention

Given that I know it has a **blue tail**,
what kind of bird is this?



Counterfactual testing

What kind of bird would this be if it
had a **green tail**?



Drawback: over-relying on label interactions




COOKING	
ROLE	VALUE
AGENT	WOMAN
FOOD	∅
HEAT	STOVE
TOOL	SPATULA
PLACE	KITCHEN

Zhao et al. 2018

Biased Model


Prone to Adversarial Examples




Gao, Lanchantin, Soffa, Qi 2018

Contributions

1. Flexibility: multi-label image classification under **arbitrary subsets** of extra or partial labels



2. Accuracy: improved results on **six different datasets** across three inference settings



3. Interactivity: state embeddings enable users to easily interact with the model and test counterfactuals




Concept Value	Class Prediction
has_yellow_underparts=1	Heermann_Gull
has_yellow_underparts=0	Glaucous_winged_Gull


Transformers can be used to exploit interactions within image scenes for accurate multi-label image classification

Conclusions and Open Problems


1. Using and discovering interactions between DNA subsequences



2. Predicting and analyzing interactions between proteins



3. Discovering and exploiting interactions between image labels




Incorporating structural interaction assumptions into the model can be extremely useful


Thesis Statement

This dissertation provides three deep learning frameworks with explicit structural assumptions to exploit, learn, predict, and understand important interactions in complex systems.


Opportunities in Genomics



Opportunities in Proteomics




Opportunities in Images and Natural Language



Context-dependent representations

Reflections


Genomics / Proteomics



ACTGACATCG

MPGKQTVKE

Images / Natural Language



this movie
is great!

Acknowledgements

Yanjun Qi



Yanjun Qi



“My feeling is that a good scientist has a great deal of what we can call curiosity. They’re just curious how things tick and they want to know the answers to questions.

- Claude Shannon


Committee



Vicente Ordoñez



Yanjun Qi



Yangfeng Ji



Clint Miller




Casey Greene

Lab-mates



Ritambhara Singh



Weilin Xu



Arshdeep Sekhon



Zeming Lin



Beilun Wang




Ji Gao



Derrick Blakely



Jack Morris



Zhe Wang



Dmitry Diochnos

Teachers, Mentors, Collaborators



Gabe Robins



Mary-Lou Soffa



Worthy Martin



Tom Weingarten



Hongning Wang




Kevin Skadron



Baoyuan Wang



Marty Humphrey



Mark Fowler



Scott Craver




Carl Betcher




Malathi Veeraraghavan

Friends




...

Brothers



Parents



Questions